


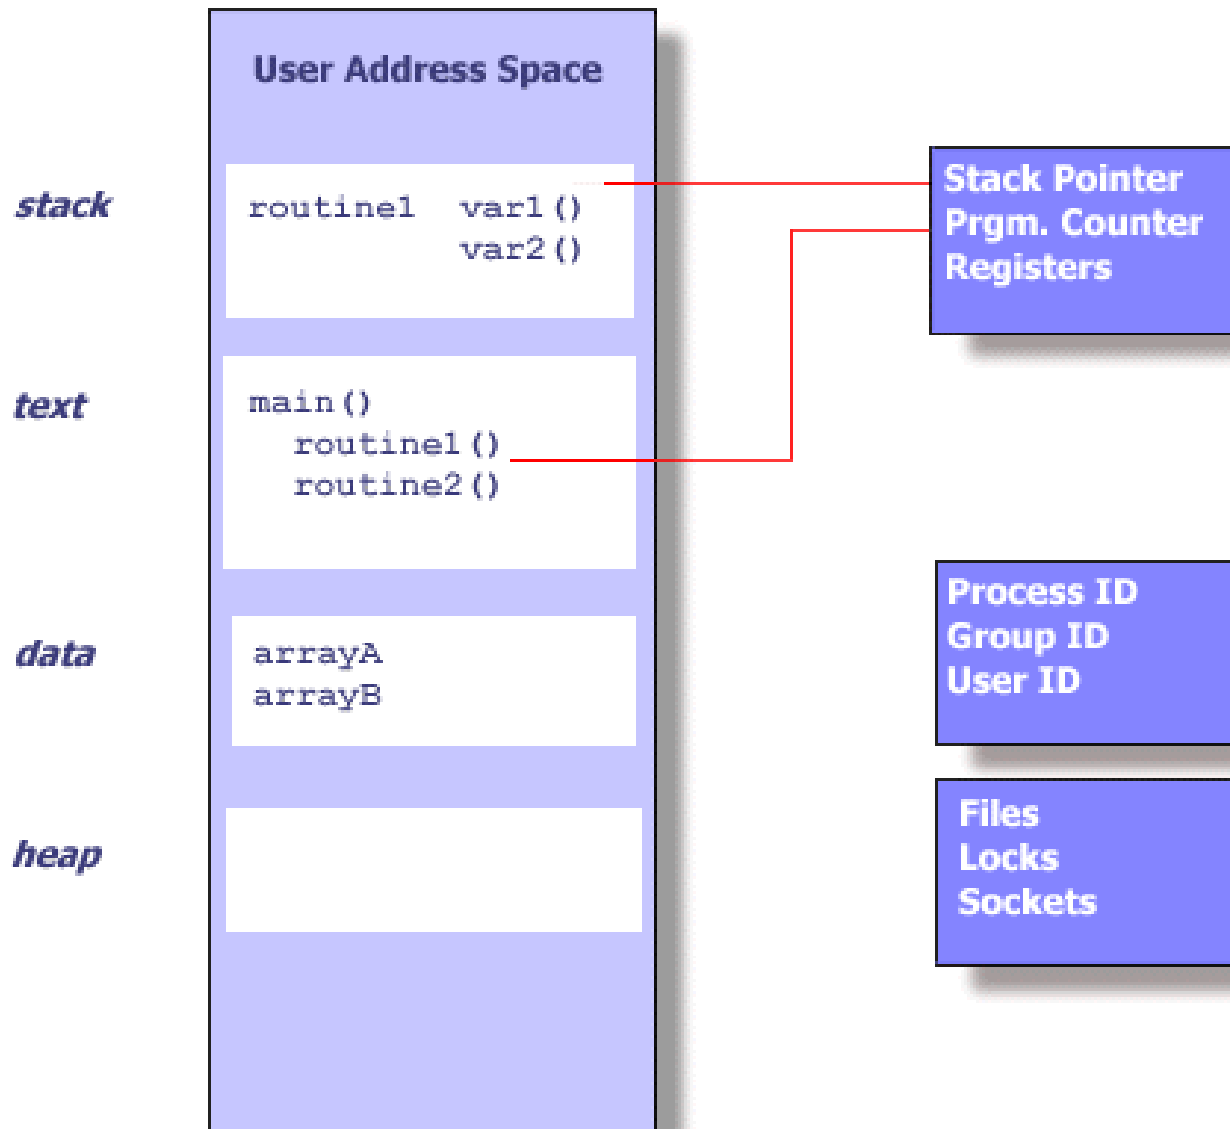
POSIX Threads Programming



What is a Thread?

- Recall that **processes** contain information about program resources and program execution state, including:
 - Process ID, process group ID, user ID, and group ID
 - Environment
 - Working directory.
 - Program instructions
 - Registers
 - Stack
 - Heap
 - File descriptors
 - Signal actions
 - Shared libraries

Unix Process

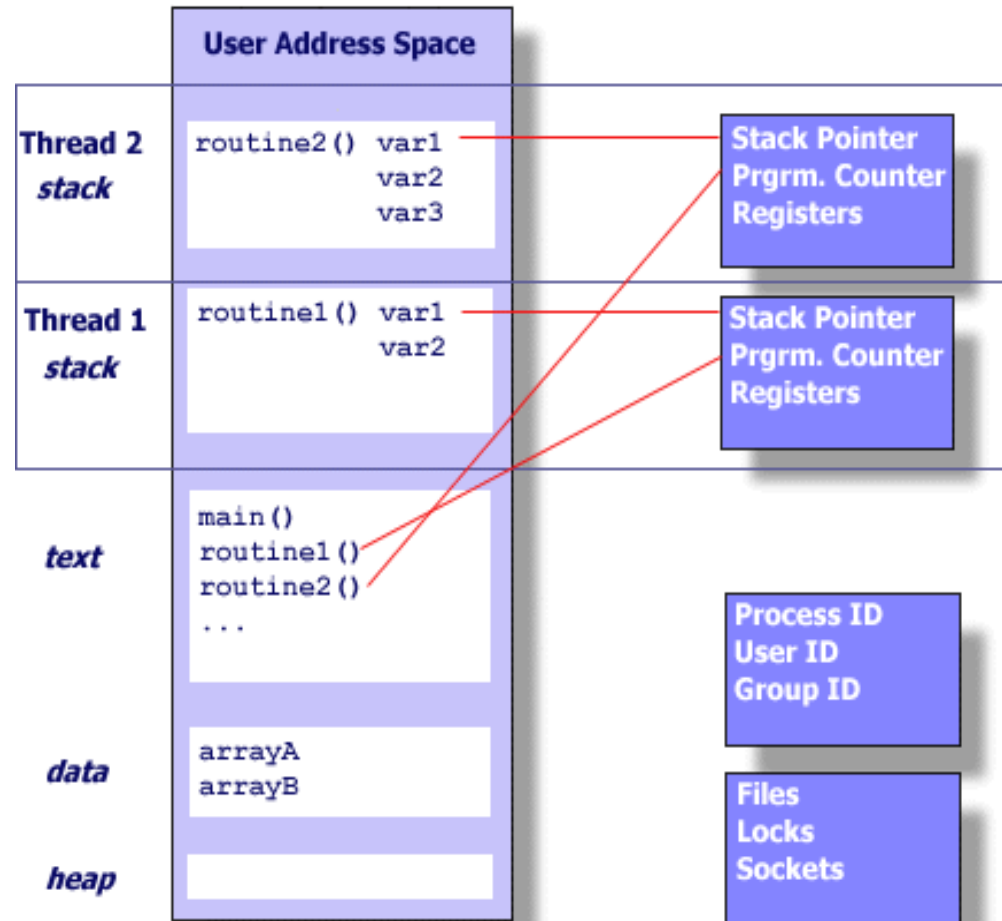


What is a Thread?

- Threads use, and exist within, the process resources
- scheduled by the operating system and run as independent entities
- duplicate only the bare essential resources that enable them to exist as executable code

What is a Thread?

- Independent flow of control possible because a thread maintains its own:
 - Stack pointer
 - Registers
 - Scheduling properties (such as policy or priority)
 - Set of pending and blocked signals
 - Thread specific data.



Summary

- In the UNIX environment a thread:
 - Exists within a process and uses the process resources
 - Has its own independent flow of control as long as its parent process exists and the OS supports it
 - Duplicates only the essential resources it needs to be independently schedulable
 - May share the process resources with other threads that act equally independently (and dependently)
 - Dies if the parent process dies - or something similar
 - Is "lightweight" because most of the overhead has already been accomplished through the creation of its process.

Thread Consequences

- Because threads within the same process share resources:
 - Changes made by one thread to shared system resources (such as closing a file) will be seen by all other threads
 - Two pointers having the same value point to the same data
 - Reading and writing to the same memory locations is possible
 - Therefore requires explicit synchronization by the programmer

Why

Pthreads

- Potential performance gains and practical advantages over non-threaded applications:
 - Overlapping CPU work with I/O
 - For example, a program may have sections where it is performing along I/O operation
 - While one thread is waiting for an I/O system call to complete, CPU intensive work can be performed by other threads.
- Priority/real-time scheduling
 - Tasks which are more important can be scheduled to supersede or interrupt lower priority tasks.
- Asynchronous event handling
 - Tasks which service events of indeterminate frequency and duration can be interleaved
 - For example, a web server can both transfer data from previous requests and manage the arrival of new requests.

Parallel Programming

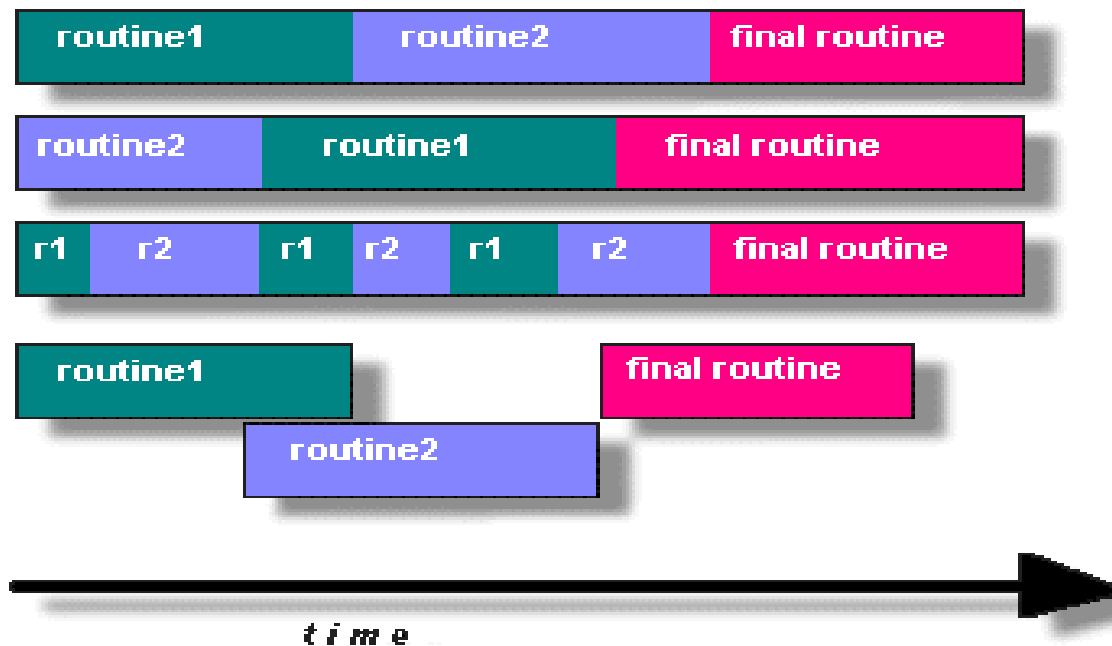
- On modern, multi-cpu machines, pthreads are ideally suited for parallel programming
- Whatever applies to parallel programming in general, applies to parallel pthreads programs

Parallel Programming

- There are many considerations for designing parallel programs, such as:
 - What type of parallel programming model to use?
 - Problem partitioning
 - Load balancing
 - Communications
 - Data dependencies
 - Synchronization and race conditions
 - Memory issues
 - I/O issues
 - Program complexity
 - Programmer effort/costs/time
 - ...

Parallel Programming

- To take advantage of Pthreads, a program must be able to be organized into discrete, independent tasks which can execute concurrently
- For example, if routine1 and routine2 can be interchanged, interleaved and/or overlapped in real time, they are candidates for threading.



Parallel Programming

- Programs having the following characteristics may be well suited for pthreads:
 - Work that can be executed, or data that can be operated on, by multiple tasks simultaneously
 - Block for potentially long I/O waits
 - Use many CPU cycles in some places but not others
 - Must respond to asynchronous events
 - Some work is more important than other work (priority interrupts)

Parallel Programming


- Pthreads can also be used for serial applications to emulate parallel execution
- For most people, the typical web browser runs on desktop/laptop machine with a single CPU
- Many things can "appear" to be happening at the same time

Parallel Programming

- Several common models for threaded programs exist:
- ***Pipeline:***
 - a task is broken into a series of sub-operations
 - each sub-operation is handled in series, but concurrently, by a different thread
 - An automobile assembly line best describes this model
- ***Peer:***
 - similar to the manager/worker model, but after the main thread creates other threads, it participates in the work.

Shared Memory Model

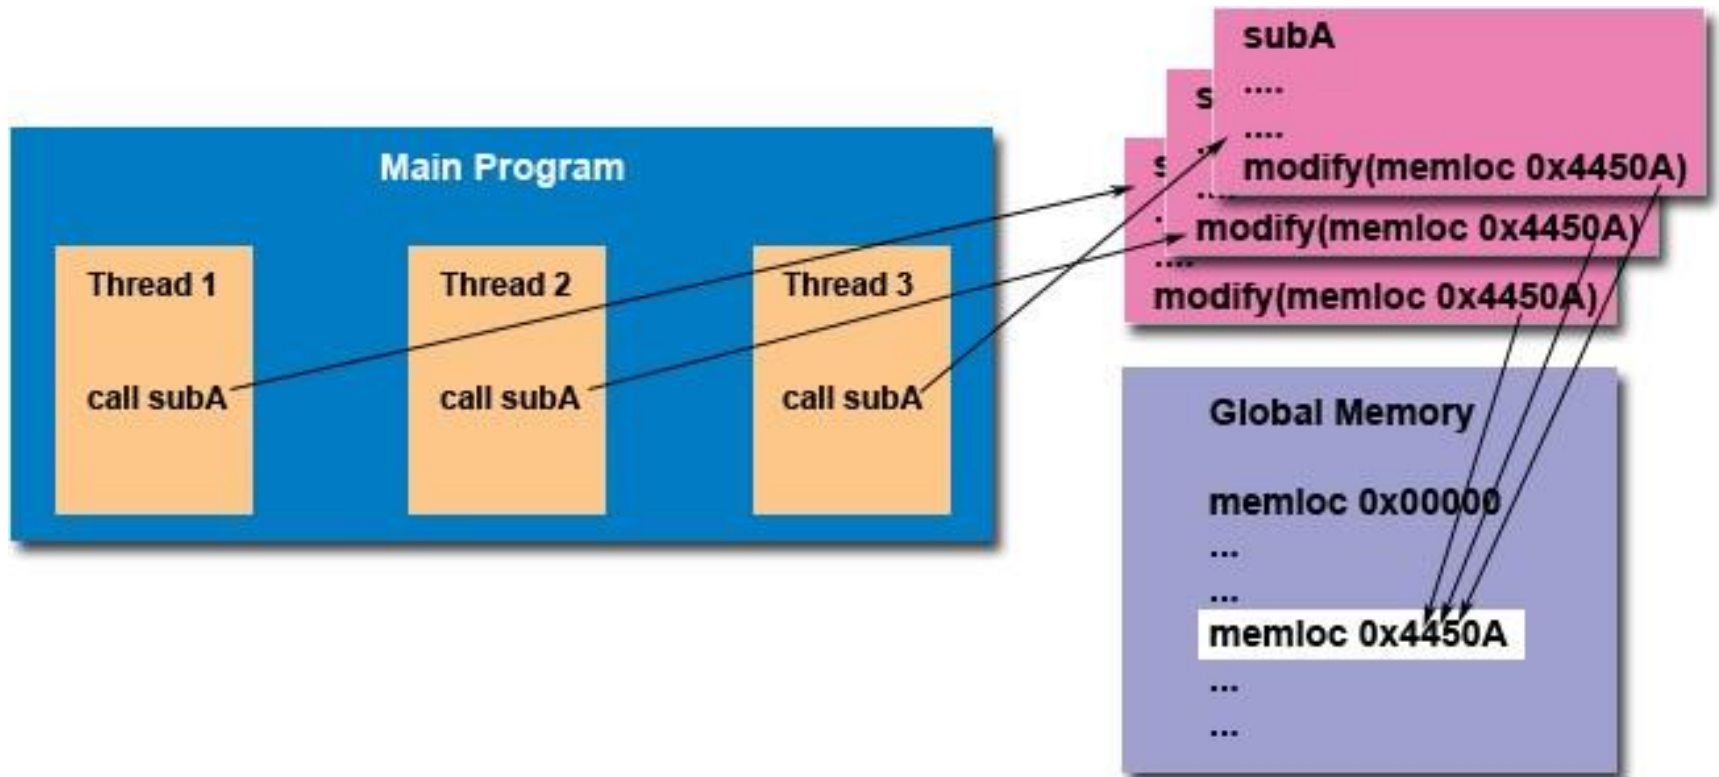
- All threads have access to the same global, shared memory
- Threads also have their own private data
- Programmers are responsible for synchronizing access (protecting) globally shared data.



Thread-safeness

- Thread-safeness: in a nutshell, refers an application's ability to execute multiple threads simultaneously without "clobbering" shared data or creating "race" conditions
- Example: an application creates several threads, each of which makes a call to the same library routine:
 - This library routine accesses/modifies a global structure or location in memory.
 - As each thread calls this routine it is possible that they may try to modify this global structure/memory location at the same time.
 - If the routine does not employ some sort of synchronization constructs to prevent data corruption, then it is not thread-safe.

Thread-safeness



Thread-safeness

- **The implication to users of external library routines:**
- if you aren't 100% certain the routine is thread-safe, then you take your chances with problems that could arise.
- **Recommendation:**
- Be careful if your application uses libraries or other objects that don't explicitly guarantee thread-safeness.
- When in doubt, assume that they are not thread-safe until proven otherwise
- This can be done by "serializing" the calls to the uncertain routine, etc.

The Pthreads API

- To compile using GNU on Linux:

gcc -pthread

- Pthread routines

pthread_create(thread, attr, start_routine, arg)

The Pthreads API

- Initially, your `main()` program comprises a single, default thread. All other threads must be explicitly created by the programmer

`pthread_create(thread, attr, start_routine, arg)`


- `pthread_create`** creates a new thread and makes it executable
- This routine can be called any number of times from anywhere within your code

Pthread_create

- `pthread_create` arguments:
 - **thread:**
An opaque, unique identifier for the new thread returned by the subroutine
 - **attr:**
An opaque attribute object that may be used to set thread attributes. You can specify a thread attributes object, or `NULL` for the default values
 - **Start_routine:**
the routine that the thread will execute once it is created
 - **arg:**
A single argument that may be passed to *start_routine*. It must be passed by reference as a pointer cast of type `void`. `NULL` may be used if no argument is to be passed.

Pthreads

- Once created, threads are peers, and may create other threads
- There is no implied hierarchy or dependency between threads



Thread Attributes

- By default, a thread is created with certain attributes
- Some of these attributes can be changed by the programmer via the thread attribute object
- `pthread_attr_init` and `pthread_attr_destroy` are used to initialize/destroy the thread attribute object
- Other routines are then used to query/set specific attributes in the thread attribute object
- Some of these attributes will be discussed later

Terminating Threads

- `pthread_exit` is used to explicitly exit a thread
- Typically, the `pthread_exit()` routine is called after a thread has completed its work and is no longer required to exist
- The programmer may optionally specify a termination *status*, which is stored as a void pointer for any thread that may join the calling thread
- Cleanup: the `pthread_exit()` routine does not close files
- Any files opened inside the thread will remain open after the thread is terminated

Example

```
#include <pthread.h> #include <stdio.h> #include  
<stdlib.h>  
#define NUM_THREADS 5  
  
void *PrintHello(void *threadid)  
{  
    long tid;  
  
    tid = (long)threadid;  
  
    printf("Hello World! It's me, thread #%ld!\n", tid);  
  
    pthread_exit(NULL);  
}
```

Example

```
int main(int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS]; int rc;
    long t;

    for(t=0;t<NUM_THREADS;t++)
    {
        printf("In main: creating thread %ld\n", t);

        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t );

        if (rc)
        {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

Example

Output:

```
In main: creating thread 0   In main: creating thread 1   In  
main: creating thread 2   In main: creating thread 3   In main:  
creating thread 4  
Hello World! It's me, thread #0!  
Hello World! It's me, thread #1!  
Hello World! It's me, thread #2!  
Hello World! It's me, thread #3!  
Hello World! It's me, thread #4!
```

Passing Arguments to Threads

- The `pthread_create()` routine permits the programmer to pass one argument to the thread start routine
- For cases where multiple arguments must be passed:
 - create a structure which contains all of the arguments
 - then pass a pointer to that structure in the `pthread_create()` routine.
 - All arguments must be passed by reference and cast to `(void*)`

Argument Passing Example 1

This code fragment demonstrates how to pass a simple integer to each thread. The calling thread uses a **unique data structure for each thread**, insuring that each thread's argument remains intact throughout the program.

```
#include <pthread.h> #include <stdio.h> #include <stdlib.h>
```

```
#define NUM_THREADS 8
```

```
char *messages[NUM_THREADS]; void *PrintHello(void *threadid)
{
    int *id_ptr, taskid;

    sleep(1);
    id_ptr = (int *) threadid; taskid = *id_ptr;
    printf("Thread %d: %s\n", taskid, messages[taskid]); pthread_exit(NULL);
}
```

Argument Passing Example 1

```
int main(int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS]; int
    *taskids[NUM_THREADS];
    int rc, t;

    messages[0] = "English: Hello World!";
    messages[1] = "French: Bonjour, le monde!";
    messages[2] = "Spanish: Hola al mundo";
    messages[3] = "Klingon: Nuq neH!";
    messages[4] = "German: Guten Tag, Welt!";
    messages[5] = "Russian: Zdravstvytye, mir!";
    messages[6] = "Japan: Sekai e konnichiwa!";
    messages[7] = "Latin: Orbis, te saluto!";
```

Argument Passing Example 1

// main, continued

```
for(t=0;t<NUM_THREADS;t++) {  
  
    taskids[t] = (int *) malloc(sizeof(int));  
    *taskids[t] = t;  
  
    printf("Creating thread %d\n", t);  
  
    rc = pthread_create(&threads[t], NULL, PrintHello, (void *) taskids[t] );  
  
    if (rc) {  
        printf("ERROR; return code from pthread_create() is %d\n", rc); exit(-  
            1);  
    }  
}  
  
pthread_exit(NULL);  
}
```

Argument Passing Example 1 Output

| | | |
|----------|--------|----------------------------|
| Creating | thread | 0 |
| Creating | thread | 1 |
| Creating | thread | 2 |
| Creating | thread | 3 |
| Creating | thread | 4 |
| Creating | thread | 5 |
| Creating | thread | 6 |
| Creating | thread | 7 |
| Thread | 0 : | English: Hello World! |
| Thread | 1 : | French: Bonjour, le monde! |
| Thread | 2 : | Spanish: Hola al mundo |

Argument Passing Example 2

```
/******  
* DESCRIPTION:  
* A "hello world" Pthreads program which demonstrates another safe way  
* to pass arguments to threads during thread creation. In this case,  
* a structure is used to pass multiple arguments.  
*****/  
  
#include <pthread.h>  
#include <stdio.h> #include <stdlib.h>  
#define NUM_THREADS 8  
char *messages[NUM_THREADS];  
struct thread_data  
{  
    int          thread_id; int    sum;  
                                char *message;  
};  
  
struct thread_data thread_data_array[NUM_THREADS];
```

Argument Passing Example 2

```
void *PrintHello(void *threadarg)
{
    int taskid, sum;  char *hello_msg;

    struct thread_data *my_data;  sleep(1);

    my_data = (struct thread_data *) threadarg;  taskid =

    my_data->thread_id;

    sum = my_data->sum;

    hello_msg = my_data->message;

    printf("Thread %d: %s    Sum=%d\n", taskid, hello_msg,
    sum);

    pthread_exit(NULL);
}
```

Argument Passing Example 2

```
int main(int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS]; int
    *taskids[NUM_THREADS];
    int rc, t, sum;

    sum=0;

    messages[0] = "English: Hello World!";
    messages[1] = "French: Bonjour, le monde!";
    messages[2] = "Spanish: Hola al mundo";
    messages[3] = "Klingon: Nuq neH!";
    messages[4] = "German: Guten Tag, Welt!";
    messages[5] = "Russian: Zdravstvyte, mir!";
    messages[6] = "Japan: Sekai e konnichiwa!";
    messages[7] = "Latin: Orbis, te saluto!";
```

Argument Passing Example 2

// main, continued

```
for(t=0;t<NUM_THREADS;t++) {  
    sum = sum + t;  
    thread_data_array[t].thread_id = t; thread_data_array[t].sum = sum;  
    thread_data_array[t].message = messages[t];  
  
    printf("Creating thread %d\n", t);  
  
    rc = pthread_create(&threads[t], NULL, PrintHello,(void *)  
thread_data_array[t] );  
    if (rc) {  
        printf("ERROR; return code from pthread_create() is %d\n", rc);  
        exit(-1);  
    }  
}  
pthread_exit(NULL);  
}
```

Argument Passing Example 2 Output

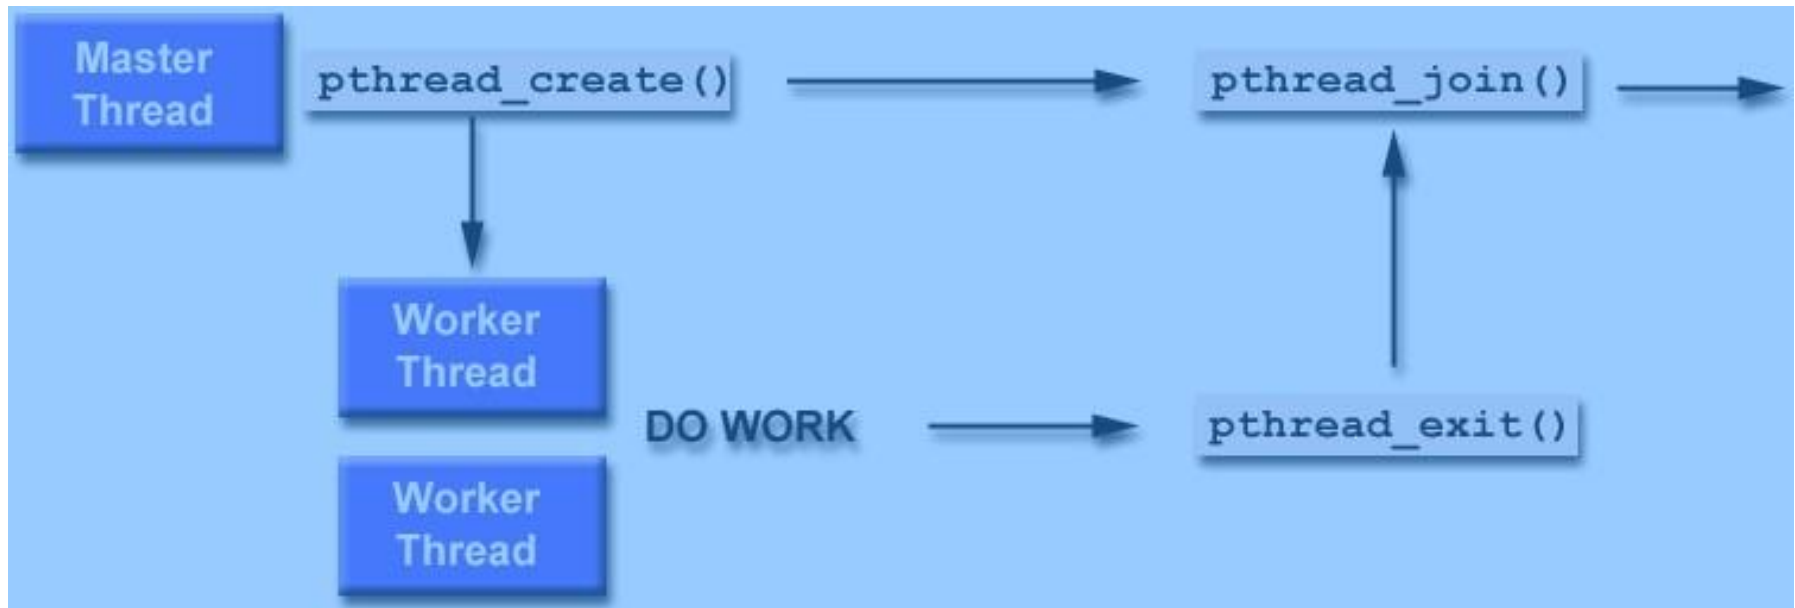
| | | |
|----------|--------|-------------------------------------|
| Creating | thread | 0 |
| Creating | thread | 1 |
| Creating | thread | 2 |
| Creating | thread | 3 |
| Creating | thread | 4 |
| Creating | thread | 5 |
| Creating | thread | 6 |
| Creating | thread | 7 |
| Thread | 0 | English: Hello World! Sum=0 |
| Thread | 1 | French: Bonjour, le monde! Sum=1 |
| Thread | 2 | Spanish: Hola al mundo Sum=3 |

Joining and Detaching Threads

- Routines:
- `pthread_join(threadid, status)`
- `pthread_detach(threadid, status)`
- `pthread_attr_setdetachstate(attr, detachstate)`
- `pthread_attr_getdetachstate(attr, detachstate)`
- The possible attribute states for the last two routines are:
 - `PTHREAD_CREATE_DETACHED` or
 - `PTHREAD_CREATE_JOINABLE`.

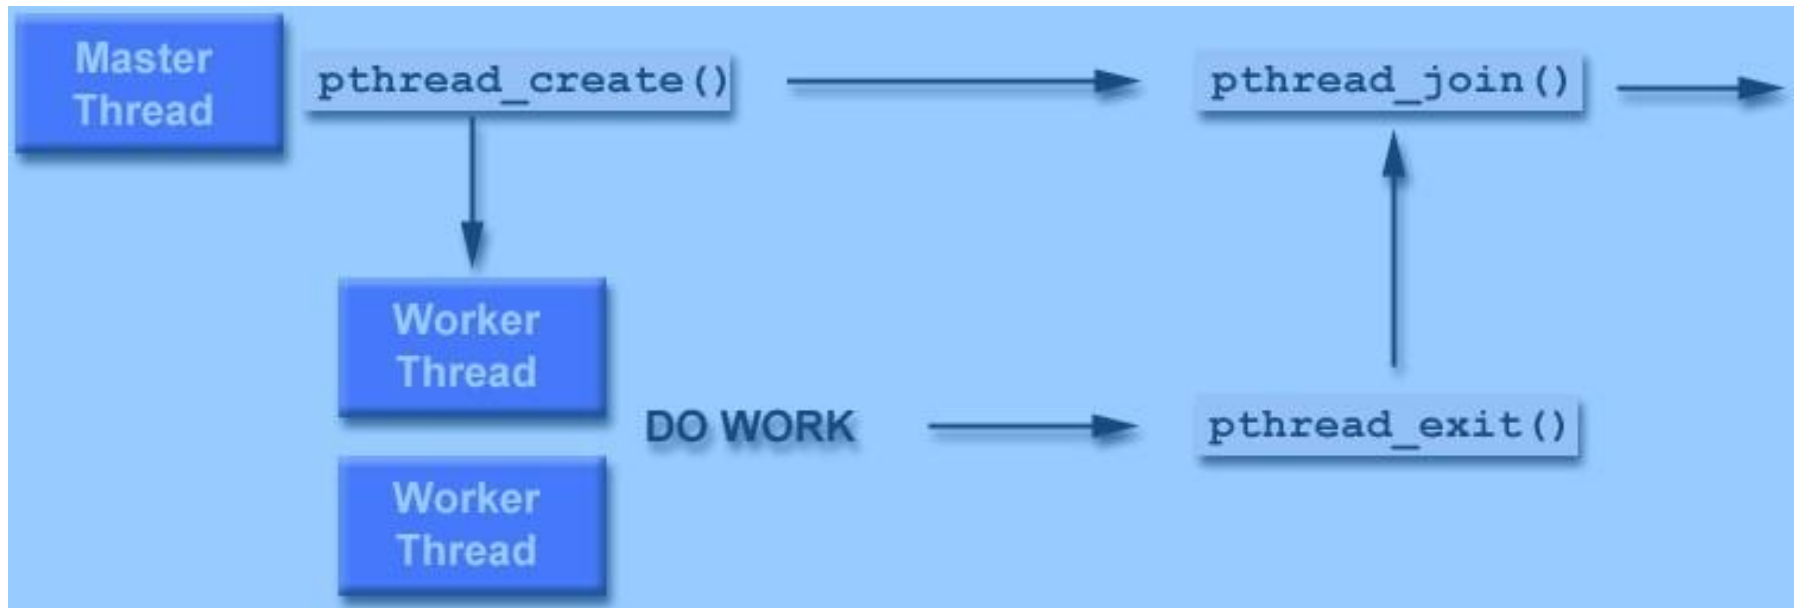
Joining

- "Joining" is one way to accomplish synchronization between threads.
- For example:



Joining

- The `pthread_join()` subroutine blocks the calling thread until the specified thread terminates
- The programmer is able to obtain the target thread's termination return status if it was specified in the target thread's call to `pthread_exit()`
- It is a logical error to attempt simultaneous multiple joins on the same target thread



Joinable or Not?

- When a thread is created, one of its attributes defines whether it is joinable or detached
- Only threads that are created as **joinable** can be joined
- If a thread is created as **detached**, it can never be joined
- The final draft of the POSIX standard specifies that threads should be created as joinable

Joinable or Not?

- To explicitly create a thread as joinable or detached, the `attr` argument in the `pthread_create()` routine is used
- **The typical 4 step process is:**
 1. Declare a pthread attribute variable of the `pthread_attr_t` data type
 2. Initialize the attribute variable with `pthread_attr_init()`
 3. Set the attribute detached status with `pthread_attr_setdetachstate()`
 4. When done, free library resources used by the attribute with `pthread_attr_destroy()`

Detach

- The `pthread_detach()` routine can be used to explicitly detach a thread even though it was created as joinable
- There is no converse routine
- **Recommendations:**
 - If a thread requires joining, consider explicitly creating it as joinable
 - This provides portability as not all implementations may create threads as joinable by default
 - If you know in advance that a thread will never need to join with another thread, consider creating it in a detached state
 - Some system resources may be able to be freed.